


sonar, sonalyze, naicreport, and the pipeline

(technical, September 2023)

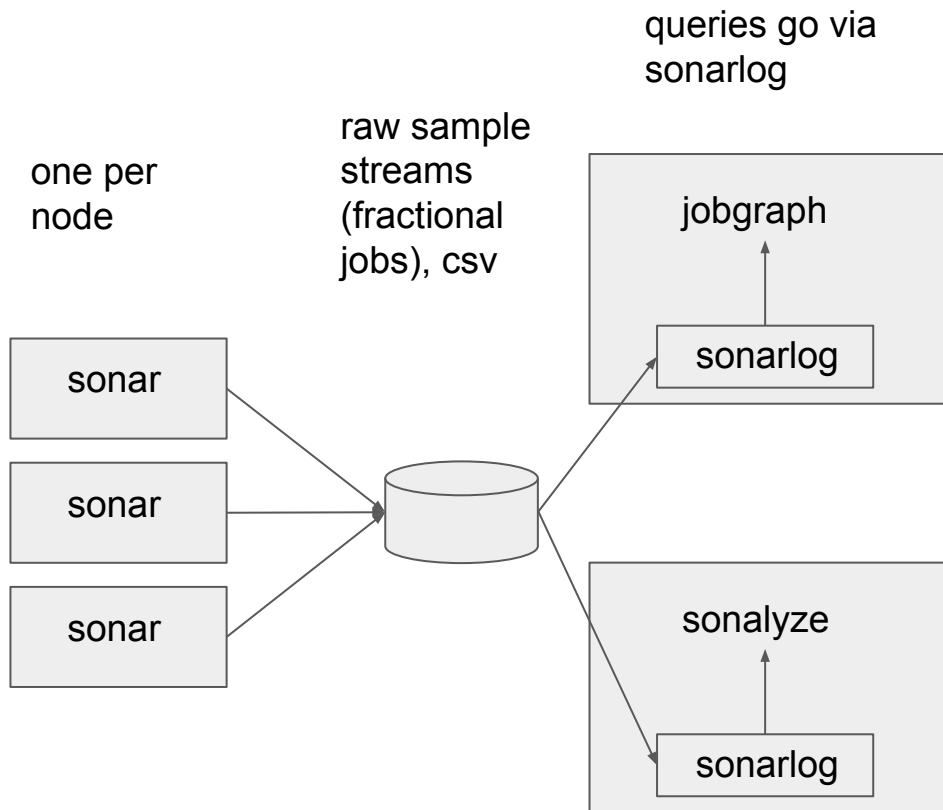
Architecture

one per
node

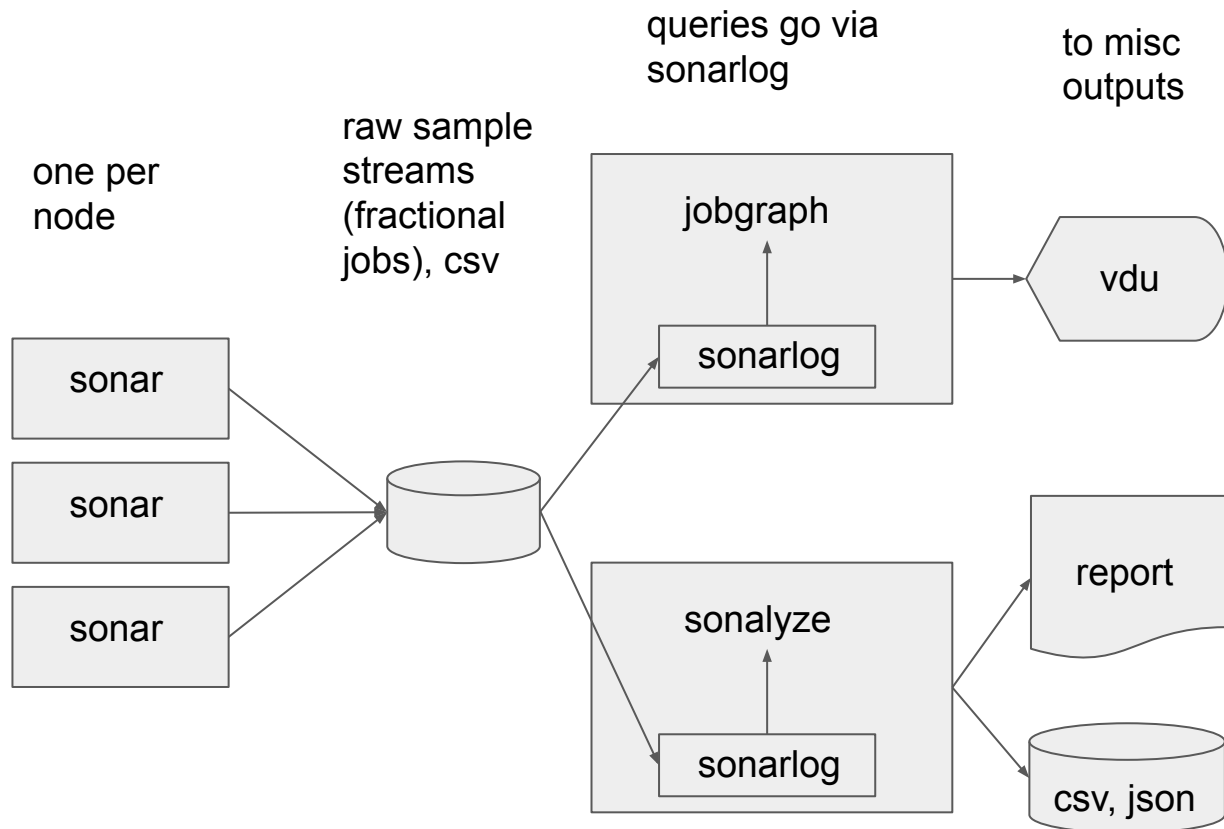
raw sample
streams
(fractional
jobs), csv



Architecture



Architecture



Here's a job report from sonalyze

```
$ ./sonalyze jobs -u- --host ml7 --data-path data
```

jobm	user	duration	host	cpu-avg	cpu-peak	mem-avg	mem-peak	gpu-avg	gpu-peak	gpumem-avg	gpumem-peak	cmd
1817618>	saeedes	0d23h55m	ml7	102	138	26	83	29	68	9	11	python3
1818070>	saeedes	0d23h55m	ml7	103	141	32	76	28	66	10	11	python3
1818403>	saeedes	0d23h55m	ml7	104	179	32	60	27	57	9	11	python3
1818623>	saeedes	0d23h55m	ml7	104	160	24	50	31	56	10	11	python3
1925362>	saeedes	0d23h55m	ml7	4	5	3	3	0	0	0	0	python3
1925587>	saeedes	0d23h55m	ml7	4	5	3	3	0	0	0	0	python3
1925829>	saeedes	0d23h55m	ml7	4	5	3	3	0	0	0	0	python3
1926449>	saeedes	0d23h55m	ml7	4	5	3	3	0	0	0	0	python3
538368	tobiasao	0d 8h20m	ml7	531	1670	6	6	2	36	1	1	python
539554	tobiasao	0d 9h10m	ml7	523	1580	6	6	2	28	1	1	python
611949>	tobiasao	0d14h30m	ml7	1575	1900	4	4	0	0	0	0	python
623616	tobiasao	0d 0h15m	ml7	24	41	6	6	4	16	1	1	python
631575>	tobiasao	0d14h25m	ml7	1665	2079	4	4	0	0	0	0	python
640122>	tobiasao	0d14h20m	ml7	1663	1918	4	4	0	0	0	0	python

There are options for selection (time, host, user, utilization, component) and formatting (almost anything)

Here's a load report, too

```
[larstha@ml4 sonar]$ ./sonalyze load --host ml7 --data-path data
```

```
HOST: ml7.hpc.uio.no
```

date	time	cpu	mem	gpu	gpumem
2023-09-21	10:00	455	123	123	37
2023-09-21	11:00	449	124	140	35
2023-09-21	12:00	449	117	146	31
2023-09-21	13:00	449	117	142	36
2023-09-21	14:00	448	115	128	26
2023-09-21	15:00	447	115	133	34
2023-09-21	16:00	453	116	125	37
2023-09-21	17:00	448	121	120	39
2023-09-21	18:00	977	128	124	40
2023-09-21	19:00	5552	139	123	41
2023-09-21	20:00	5876	137	102	40
2023-09-21	21:00	5894	139	104	35
2023-09-21	22:00	5803	140	110	37
2023-09-21	23:00	5848	146	86	37
2023-09-22	00:00	5860	149	89	33
2023-09-22	01:00	5858	148	91	29
2023-09-22	02:00	5897	152	73	31
2023-09-22	03:00	5826	151	102	36
2023-09-22	04:00	5829	145	113	36
2023-09-22	05:00	5847	137	115	38
2023-09-22	06:00	5847	140	100	40
2023-09-22	07:00	5841	138	110	36
2023-09-22	08:00	5840	136	103	33
2023-09-22	09:00	5838	140	113	34
2023-09-22	10:00	5841	138	112	33

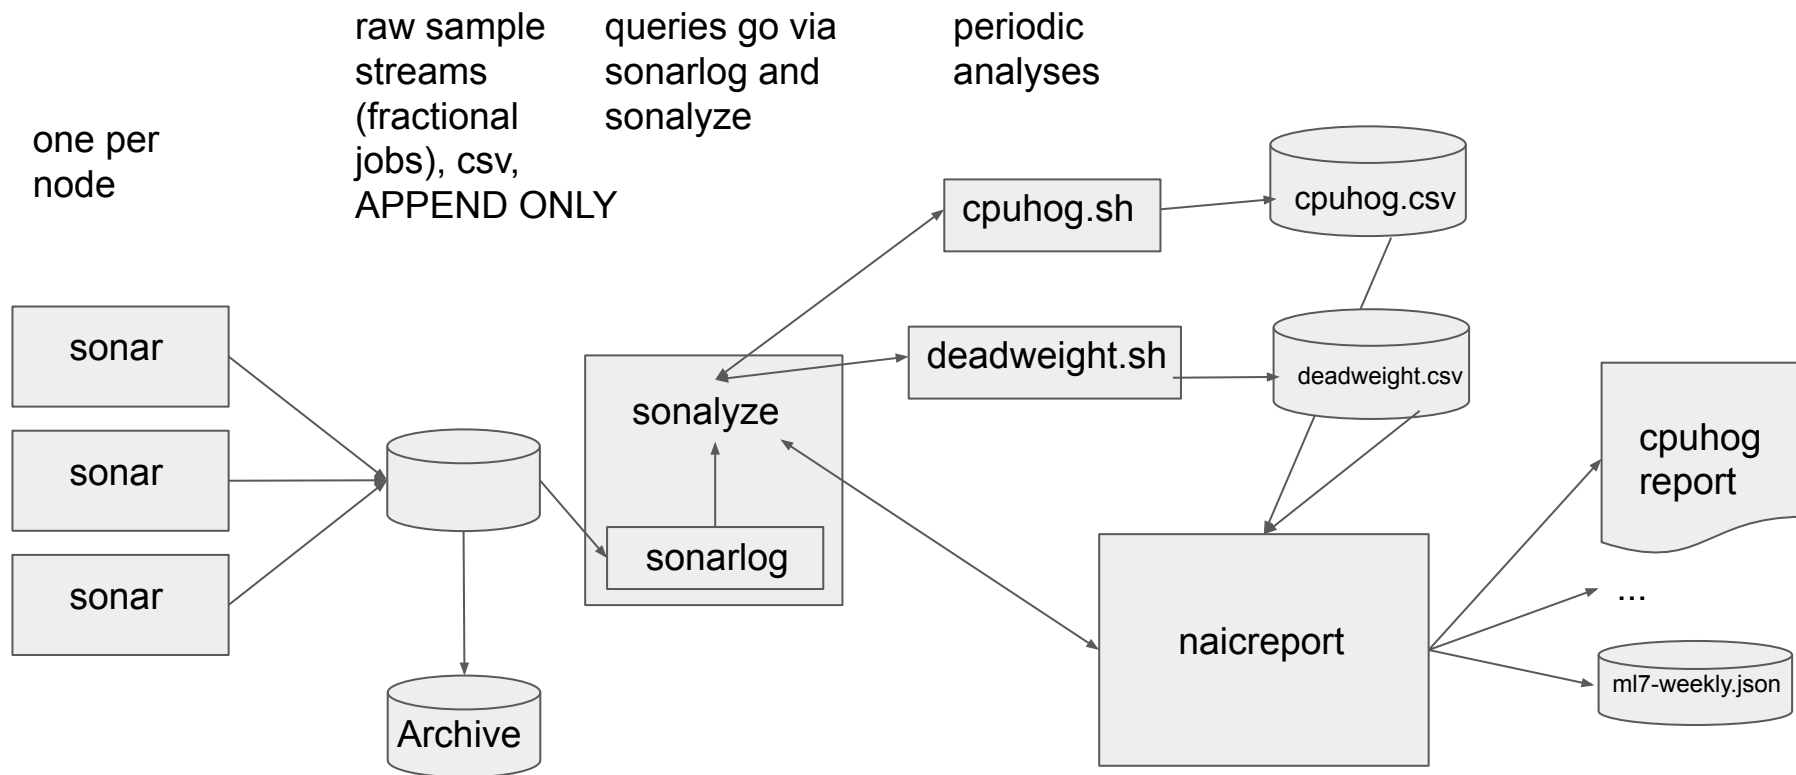
Load averaged hourly here

Daily also available

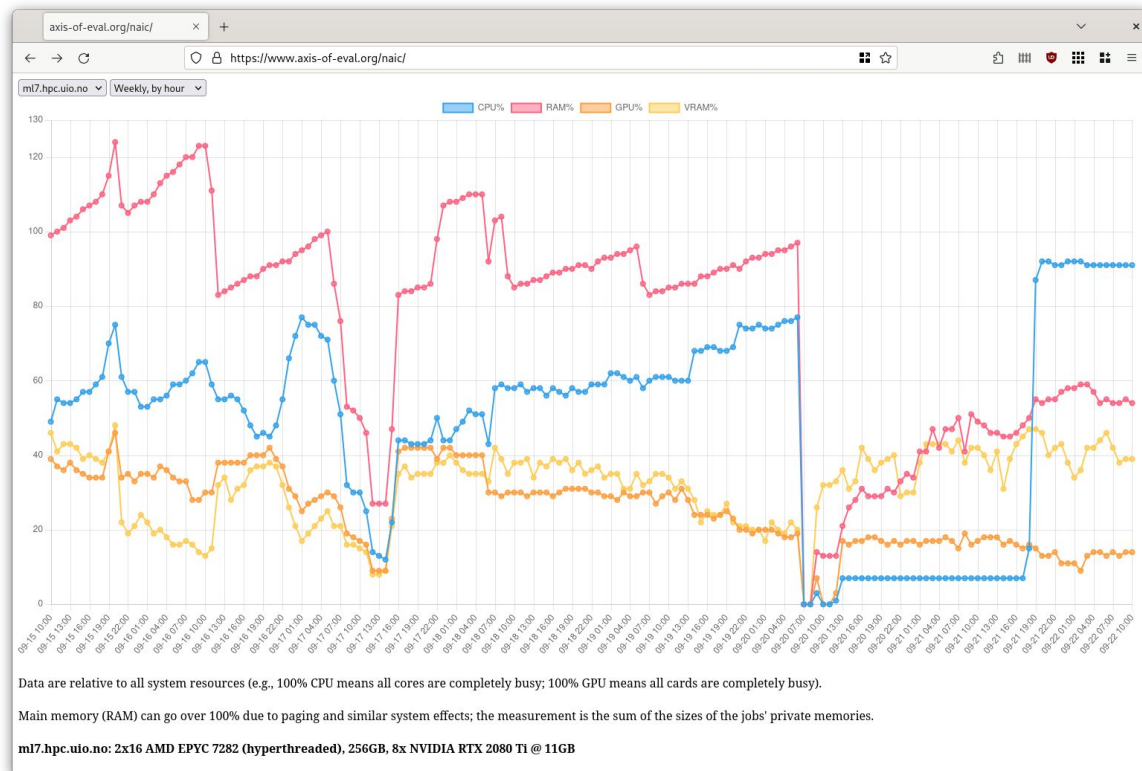
Many selection options

Many formatting options (also csv, json)

Architecture



Rendered view of ml7-weekly.json (load report)



Cpuhog report (excerpt)

New CPU hog detected (uses a lot of CPU and no GPU) on host "ml6":

Job#: 2712710

User: hermanno

Command: kited

Started on or before: 2023-09-07 10:00

Violation first detected: 2023-09-22 11:13

Observed data:

CPU peak = 37 cores

CPU utilization avg/peak = 1%, 58%

Memory utilization avg/peak = 5%, 6%

New CPU hog detected (uses a lot of CPU and no GPU) on host "ml6":

Job#: 3043187

User: poyenyt

Command: python3.9

Started on or before: 2023-09-07 07:50

Violation first detected: 2023-09-22 11:13

Observed data:

CPU peak = 15 cores

CPU utilization avg/peak = 5%, 24%

Memory utilization avg/peak = 3%, 3%

Summary wall of text

- *sonar* runs on every node (by cron), fairly often (every 5 minutes on ml systems, about 2MB uncompressed text per day across all nodes), generates raw sample streams in csv form
- *sonarlog* ingests sonar data and cleans them up for general utility:
 - add missing data
 - fixup questionable fields
 - filter records and jobs according to input arguments
 - segregate all sample streams by the job artifact key (host, job-id, cmd)
 - provide utilities to merge streams predictably and to correctly build aggregates. somewhat tricky.
 - sonarlog is pretty clean now and can probably be reused from other programs (eg jobgraph)
- *sonalyze* operates on the cleaned-up sample streams
 - aggregates data for jobs and hosts and prints these aggregates (multi-node jobs not 100% done probably)
 - command line switches for selecting input records, aggregating in different ways, printing in different ways
- cron jobs run *sonalyze* every 2 hrs for *cpuhog* and *deadweight* analysis, generates more csv
- *sonalyze* can also be run manually for ad-hoc queries, to better understand summarized data
- *naicreport* is run by cron occasionally to generate reports for the web front end (load, cpuhog, deadweight)
- *naicreport* can also be run manually as needed
- currently this pipeline is only on the ML nodes
- *sonar*, *sonarlog* and *sonalyze* written in Rust; *naicreport* written in Go; bash to glue everything together; HTML+JS for Web
- <https://github.com/NAICNO/Jobanalyzer>

Status + Future work

- sonar and sonarlog are very stable now, minor bug + feature work
 - definitely want to add some logging of communication work
- sonarlog is definitely shareable with jobgraph, probably a good idea
- sonalyze mostly stable but will see some more use-case driven work
- naicreport is growing daily, will continue with this
- NAIC
 - production deployment on ML nodes + hosted analysis & web is imminent
 - multi-node work: Fox, then the world!
 - flesh out NAIC use cases and requirements in more detail, then address them
 - probably a lot of experimentation to see what works
- data management:
 - compress and archive older data to keep the volume under control
 - allow sonarlog to work on the compressed data transparently: older data sometimes useful, long-term reports
- presentation
 - we have a UI intern, might try to make it look nice & more functional
 - lots of reports driven by NAIC use cases, TBD
- (more stuff i haven't thought about, most of it appears in the bug tracker anyhow)